

Linear Regression

Modelo: $p(y|\mathbf{x}, \mathbf{w}) = \mathcal{N}(y|\mathbf{w}^\top \mathbf{x}, \sigma^2)$
Predicción: $\hat{y} = \mathbb{E}[y|\mathbf{x}, \mathbf{w}] = \mathbf{w}^\top \mathbf{x}$
 $\mathcal{L} = \text{MSE} = \frac{1}{N} \sum_i (y_i - \mathbf{w}^\top \mathbf{x}_i)^2$
 $\nabla_{\mathbf{w}} \mathcal{L} = \frac{2}{N} \sum_i (\mathbf{w}^\top \mathbf{x}_i - y_i) \mathbf{x}_i = \frac{2}{N} \mathbf{X}^\top (\mathbf{X} \mathbf{w} - \mathbf{y})$
Solución (MLE / LS): $\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$

Multi-layer Perceptron (MLP)

$\mathbf{z}^{(\ell)} \in \mathbb{R}^{M^{(\ell)}}:$ output capa ℓ
 $\mathbf{a}^{(\ell)} \in \mathbb{R}^{M^{(\ell)}}:$ preactivación capa ℓ
 $\delta^{(\ell)} \in \mathbb{R}^{M^{(\ell)}}:$ error de propagación en capa ℓ
 $\mathbf{W}^{(\ell)} \in \mathbb{R}^{M^{(\ell)} \times M^{(\ell-1)}}:$ params entre ℓ y $\ell-1$
Back-propagation
Forward-pass
 $\mathbf{z}^{(0)} \leftarrow \mathbf{x}_i$
for $\ell = 1 : L$ do
 $\mathbf{a}^{(\ell)} \leftarrow \mathbf{W}^{(\ell)} \mathbf{z}^{(\ell-1)} + \mathbf{b}^{(\ell)}$
 $\mathbf{z}^{(\ell)} \leftarrow \varphi^{(\ell)}(\mathbf{a}^{(\ell)})$
 $\hat{\mathbf{y}}_i \leftarrow \mathbf{z}^{(L)}, \quad \mathcal{L}_i \leftarrow \mathcal{L}_i(\mathbf{y}_i, \hat{\mathbf{y}}_i)$
Backward-pass
 $\delta^{(L)} \leftarrow \varphi'^{(L)}(\mathbf{a}^{(L)}) \odot \frac{\partial \mathcal{L}_i}{\partial \hat{\mathbf{y}}_i}$
for $l = L-1 : 1$ do
 $\delta^{(l)} \leftarrow \varphi'^{(l)}(\mathbf{a}^{(l)}) \odot \mathbf{W}^{(l+1)\top} \delta^{(l+1)}$
 $\frac{\partial \mathcal{L}_i}{\partial \mathbf{W}^{(l)}} \leftarrow \delta^{(l)} \mathbf{z}^{(l-1)\top}$
 $\frac{\partial \mathcal{L}_i}{\partial \mathbf{b}^{(l)}} \leftarrow \delta^{(l)}$

Activation Functions (φ)

Sigmoid: $\sigma(z) = \frac{1}{1+e^{-z}}; \quad \sigma' = \sigma(1-\sigma)$
Tanh: $\tanh' = 1 - \tanh^2$
ReLU: $\max(0, z); \quad \text{ReLU}' = \mathbb{I}[z > 0]$
Leaky ReLU: $\max(\alpha z, z), \quad \alpha \ll 1$
Softplus: $\log(1 + e^z)$
Identity: $\varphi(z) = z$
Softmax: $\sigma(\mathbf{z})_i = e^{z_i} / \sum_j e^{z_j}$

Ensemble Methods

Stacking: combina modelos heterogeneos
Bagging: paralelo, reduce varianza
• **Random Forest:** + feature bagging.
Boosting: secuencial, reduce bias
• **Grad. Boosting** (ajusta residuos con GD)
• XGBoost, LightGBM, CatBoost
• **AdaBoost** (ponderación mayor a muestras mal clasificadas)

Distributions + MLE estimators

Gaussian $\mathcal{N}(\mathbf{x}|\mu, \sigma^2):$
$$\frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

 $\hat{\mu} = \frac{1}{N} \sum_i x_i \quad \hat{\sigma}^2 = \frac{1}{N} \sum_i (x_i - \hat{\mu})^2$
Gaussian Multivariada $\mathcal{N}(\mathbf{x}|\mu, \Sigma):$
$$\frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}} \exp\left(-\frac{1}{2} (\mathbf{x} - \mu)^\top \Sigma^{-1} (\mathbf{x} - \mu)\right)$$

 $\hat{\mu} = \frac{1}{N} \sum_i \mathbf{x}_i \quad \hat{\Sigma} = \frac{1}{N} \sum_i (\mathbf{x}_i - \hat{\mu})(\mathbf{x}_i - \hat{\mu})^\top$
Bernoulli $\text{Ber}(x|\mu):$
$$\mu^x (1-\mu)^{1-x}$$

 $\hat{\mu} = \frac{1}{N} \sum_i x_i$
Catagórica $\text{Cat}(y|\pi):$
$$\prod_c \pi_c^{\mathbb{I}[y=c]}, \quad \sum_c \pi_c = 1$$

 $\hat{\pi}_c = \frac{N_c}{N}$ (frecuencia de clase c)
Laplace $\text{Lap}(x|\mu, b):$
$$\frac{1}{2b} \exp\left(-\frac{|x-\mu|}{b}\right)$$

 $\hat{\mu} = \text{median}(\{x_i\}) \quad \hat{b} = \frac{1}{N} \sum_i |x_i - \hat{\mu}|$
Binomial $\text{Binom}(k|n, p):$
$$\binom{n}{k} p^k (1-p)^{n-k}$$

 $\hat{p} = \frac{1}{N} \sum_i \frac{k_i}{n}$

Logistic Regression

Binaria
Modelo: $p(y|\mathbf{x}, \boldsymbol{\theta}) = \text{Ber}(y|\sigma(\mathbf{w}^\top \mathbf{x} + b))$
Predicción (prob): $\mu_i = \sigma(\mathbf{w}^\top \mathbf{x}_i)$
 $\mathcal{L} = \frac{1}{N} \sum_i \mathbb{H}_{ce}(y_i, \mu_i)$
 $\nabla_{\mathbf{w}} \mathcal{L} = \frac{1}{N} \sum_i (\mu_i - y_i) \mathbf{x}_i = \frac{1}{N} \mathbf{X}^\top (\boldsymbol{\mu} - \mathbf{y})$
Multiclase
Modelo: $p(y|\mathbf{x}, \mathbf{w}) = \text{Cat}(y|\text{Softmax}(\mathbf{w}^\top \mathbf{x}))$
Predicción (prob): $\boldsymbol{\mu}_i = \text{Softmax}(\mathbf{w}^\top \mathbf{x}_i)$
 $\mathcal{L} = \frac{1}{N} \sum_i \mathbb{H}_{ce}(\mathbf{y}_i, \boldsymbol{\mu}_i)$
 $\nabla_{\mathbf{w}} \mathcal{L} = \frac{1}{N} \sum_i \mathbf{x}_i (\boldsymbol{\mu}_i - \mathbf{y}_i)^\top = \frac{1}{N} \mathbf{X}^\top (\mathbf{M} - \mathbf{Y})$
($\mathbf{M}$ matriz de predicciones, $\mathbf{Y}$ one-hots)

Naive Bayes

Naive Bayes Assumption: $x_d \perp x_{d'} \mid y = c$
Modelo: $p(y = c|\mathbf{x}) \propto \pi_c \prod_{d=1}^D p(x_d|y = c, \boldsymbol{\theta}_{dc})$
 $p(x_d|y = c, \boldsymbol{\theta}_{dc})$ según tipo de x_d :
• Binaria $\text{Ber}(x_d|\theta_{dc})$
• Categórica = $\text{Cat}(x_d|\boldsymbol{\theta}_{dc})$
• Real = $\mathcal{N}(x_d|\mu_{dc}, \sigma_{dc}^2)$
Entrenamiento: MLE

Non-Parametric Methods

KNN-classifier:
 $p(y = c|\mathbf{x}, \mathcal{D}) = \frac{1}{K} \sum_{i \in \mathcal{N}_K(\mathbf{x}, \mathcal{D})} \mathbb{I}(y_i = c)$
Parzen Window Density Estimator:
 $p(\mathbf{x}, \mathcal{D}) = \frac{1}{N} \sum_{i=1}^N \mathcal{K}_h(\mathbf{x}, \mathbf{x}_i)$

NN Regularization + ∇ vanishing sol.

Weight Decay (ℓ_2): MAP estimation
 $p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{0}, \alpha^2 \mathbf{I}), p(\mathbf{b}) = \mathcal{N}(\mathbf{b}|\mathbf{0}, \beta^2 \mathbf{I})$
Dropout: $\theta_{lji} = w_{lji} \epsilon_{li}, \epsilon \sim \text{Ber}(1-p)$
Al predecir: escalar activaciones por $(1-p)$
Early Stopping: frenar cuando $\uparrow \mathcal{L}_{\text{val}}$
Batch Normalization: normalizar activaciones por mini-batch $\mathcal{B}$
$$\hat{z}^{(\ell)} = \frac{z^{(\ell)} - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}}, \quad \tilde{z}^{(\ell)} = \gamma \hat{z}^{(\ell)} + \beta$$

Glorot/Xavier initialization
 $w_{ij} \sim \mathcal{N}(0, \sigma^2)$ con $\sigma^2 = 2/(n_{\text{in}} + n_{\text{out}})$
Activaciones sin saturación: ReLU, Leaky ReLU (evitan saturación)
Gradient clipping: $\mathbf{g} = \min\left(0, \frac{c}{\|\mathbf{g}\|}\right) \mathbf{g}$
Batch/Layer Normalization

Convexity

$f(\lambda \mathbf{x} + (1-\lambda) \mathbf{y}) \leq \lambda f(\mathbf{x}) + (1-\lambda) f(\mathbf{y})$
1er orden: $f(\mathbf{y}) \geq f(\mathbf{x}) + \nabla f(\mathbf{x})^\top (\mathbf{y} - \mathbf{x})$
2do orden: $\nabla^2 f(\mathbf{x}) \succeq 0 \forall \mathbf{x}$ (Hessiano SDP)

Maximum A Posteriori (MAP)

$\mathbf{w}^* = \arg \max_{\mathbf{w}} p(\mathbf{w}|\mathcal{D}) = \arg \max_{\mathbf{w}} p(\mathcal{D}|\mathbf{w}) p(\mathbf{w})$
 $\ell_2 \rightarrow p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{0}, \lambda^{-1} \mathbf{I}) \rightarrow \mathcal{L} + \lambda \|\mathbf{w}\|^2$
 $\ell_1 \rightarrow p(\mathbf{w}) = \text{Laplace}(\mathbf{w}|\mathbf{0}, \lambda^{-1}) \rightarrow \mathcal{L} + \lambda \|\mathbf{w}\|_1$

Bayes

$$p(\mathbf{w}|\mathcal{D}) = \frac{p(\mathcal{D}|\mathbf{w}) p(\mathbf{w})}{p(\mathcal{D})} \propto p(\mathcal{D}|\mathbf{w}) p(\mathbf{w})$$

$$\text{posterior} \propto \text{likelihood} \times \text{prior}$$

Information Theory

Entropy: $\mathbb{H}(\mathbf{p}) = -\sum_i p_i \log p_i$
Cross-E: $\mathbb{H}_{ce}(\mathbf{p}, \mathbf{q}) = -\sum_i p_i \log q_i$
Binary Cross-E:
 $\mathbb{H}_{ce}(p, q) = -[p \log q + (1-p) \log(1-q)]$
Joint-E: $\mathbb{H}(X, Y) = -\sum p(x, y) \log p(x, y)$
Conditional-E: $\mathbb{H}(Y|X) = \mathbb{H}(X, Y) - \mathbb{H}(X)$
Relative-E (KL divergence)
 $\mathbb{H}(p||q) = \sum_i p_i \log \frac{p_i}{q_i} = \mathbb{H}(p||q) - \mathbb{H}(p)$
Mutual Information
 $I(X; Y) = \mathbb{H}(p(x, y) || p(x)p(y))$
 $I(X; Y) = 0 \Leftrightarrow X \perp Y$

GDA

Modelo: $p(y = c|\mathbf{x}, \mathbf{w}) \propto \pi_c p(\mathbf{x}|y = c, \mathbf{w})$
Class-cond: $p(\mathbf{x}|y = c, \mathbf{w}) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_c, \boldsymbol{\Sigma}_c)$
Prior: $\pi_c = p(y = c|\mathbf{w})$
Entrenamiento: MLE ($\hat{\pi}, \hat{\boldsymbol{\mu}}_c, \hat{\boldsymbol{\Sigma}}_c$)
MAP: $\hat{\mathbf{S}}_{\text{map}} = \lambda \text{diag}(\hat{\boldsymbol{\Sigma}}) + (1-\lambda) \hat{\boldsymbol{\Sigma}}$
QDA ($\boldsymbol{\Sigma}_c$ libre por clase)
Frontera: $\mathbf{x}^\top \mathbf{A} \mathbf{x} + \mathbf{b}^\top \mathbf{x} + c_0 = 0$
 $\mathbf{A} = \frac{1}{2} (\boldsymbol{\Sigma}_{c'}^{-1} - \boldsymbol{\Sigma}_c^{-1}), \mathbf{b} = \boldsymbol{\Sigma}_c^{-1} \boldsymbol{\mu}_c - \boldsymbol{\Sigma}_{c'}^{-1} \boldsymbol{\mu}_{c'}$
 $c_0 = \log \frac{\pi_c}{\pi_{c'}} - \frac{1}{2} \log \frac{|\boldsymbol{\Sigma}_c|}{|\boldsymbol{\Sigma}_{c'}|} - \frac{1}{2} (\boldsymbol{\mu}_c^\top \boldsymbol{\Sigma}_c^{-1} \boldsymbol{\mu}_c - \boldsymbol{\mu}_{c'}^\top \boldsymbol{\Sigma}_{c'}^{-1} \boldsymbol{\mu}_{c'})$
LDA ($\boldsymbol{\Sigma}_c = \boldsymbol{\Sigma} \forall c$)
Frontera: $\mathbf{w}^\top \mathbf{x} + w_0 = 0$
donde $\mathbf{w} = \boldsymbol{\Sigma}^{-1} (\boldsymbol{\mu}_c - \boldsymbol{\mu}_{c'})$
 $w_0 = \log \frac{\pi_c}{\pi_{c'}} - \frac{1}{2} (\boldsymbol{\mu}_c + \boldsymbol{\mu}_{c'})^\top \mathbf{w}$

CART

$\mathcal{L}(\boldsymbol{\theta}) = \sum_i \ell(y_i, f(\mathbf{x}_i; \boldsymbol{\theta}))$
Split (nodo i):
$$\arg \min_j \min_t \frac{|\mathcal{D}_i^L|}{|\mathcal{D}_i|} C^L(\mathcal{D}_i^L) + \frac{|\mathcal{D}_i^R|}{|\mathcal{D}_i|} C^R(\mathcal{D}_i^R)$$

Regresión: $C(\mathcal{D}_i) = \text{MSE}$
Clasificación: $C(\mathcal{D}_i) = G_i = \sum_c \hat{\pi}_{ic} (1 - \hat{\pi}_{ic})$ ó
 $C(\mathcal{D}_i) = H_i = \mathbb{H}(\hat{\boldsymbol{\pi}}_i) = -\sum_{ic} \hat{\pi}_{ic} \log(\hat{\pi}_{ic})$

Perceptron

$\hat{y} = \mathbb{I}(\mathbf{w}^\top \mathbf{x} + b > 0) \rightarrow$ (determinístico)
Update: $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - \eta (\hat{y}_i - y_i) \mathbf{x}_i$

Optimization

GD: $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - \eta_t \nabla_{\mathbf{w}} \mathcal{L} |_{\mathbf{w}_t}$
SGD: GD con una muestra $(\mathbf{x}_i, y_i)$ por update
Mini-batch: $\mathcal{B} \subset \mathcal{D}$ por update
Momentum: acumula gradientes pasados
 $\mathbf{m}_t \leftarrow \beta \mathbf{m}_{t-1} + \nabla \mathcal{L} |_{\mathbf{w}_t}, \quad \mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - \eta_t \mathbf{m}_t$
Newton: $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - \eta_t \mathbf{H}_t^{-1} \nabla_{\mathbf{w}} \mathcal{L}$
LR Schedule: $\downarrow \eta_t$ por iteración
• Lineal: $\eta_t = (1 - \frac{t}{T}) \eta_0 + (\frac{t}{T}) \eta_T$
• Potencia: $\eta_t = \eta_0 (1 + \frac{t}{s})^c, \quad c, s$ params
• Exponencial: $\eta_t = \eta_0 \cdot c^{(t/s)}, \quad c, s$ params
AdaGrad: η_t adaptativo por parámetro
Adam: Momentum + AdaGrad (η y m por parámetro)

Linear Algebra

SDP: $\mathbf{A} \succeq 0 \Leftrightarrow \mathbf{x}^\top \mathbf{A} \mathbf{x} \geq 0 \forall \mathbf{x}$.
 $\mathbf{A}^\top \mathbf{A}, \mathbf{A} \mathbf{A}^\top \succeq 0$ siempre
Cond: $\kappa(\mathbf{A}) = \sigma_{\max}/\sigma_{\min}, \kappa(\mathbf{A}^\top \mathbf{A}) = \kappa(\mathbf{A})^2$
Precisión: $\boldsymbol{\Delta} \triangleq \boldsymbol{\Sigma}^{-1}$
Derivadas conocidas
• $\nabla_{\mathbf{x}}(\mathbf{x}^\top \mathbf{A} \mathbf{x}) = 2\mathbf{A} \mathbf{x}$ ($\mathbf{A}$ simétrica)
• $\nabla_{\mathbf{x}}(\mathbf{a}^\top \mathbf{x}) = \mathbf{a}$
• $\nabla_{\mathbf{S}} \log |\mathbf{S}| = \boldsymbol{\Sigma}^{-\top}$
Traza: $\mathbf{x}^\top \mathbf{A} \mathbf{x} = \text{tr}(\mathbf{A} \mathbf{x} \mathbf{x}^\top), \text{tr}(\mathbf{A} \mathbf{B}) = \text{tr}(\mathbf{B} \mathbf{A})$
Jacobiano: $\mathbf{J} = \frac{\partial \mathbf{f}}{\partial \mathbf{x}} \in \mathbb{R}^{m \times n}, J_{ij} = \frac{\partial f_i}{\partial x_j}$
Si $m = 1 \implies \mathbf{J} = \nabla f^\top$.

Regression Metrics

$e_i \triangleq y_i - \hat{y}_i; \text{RSS} \triangleq \sum_i e_i^2; \text{TSS} \triangleq \sum_i (y_i - \bar{y})^2$
 $\text{MSE} = \frac{1}{N} \text{RSS}$
 $\text{RMSE} = \sqrt{\text{MSE}}$
 $\text{MAE} = \frac{1}{N} \sum_i |e_i|$
 $R^2 \triangleq 1 - \frac{\text{RSS}}{\text{TSS}} \in [0, 1]$ (varianza explicada)

Classification Metrics

$\text{Acc} = \frac{TP+TN}{TP+TN+FP+FN} \quad \text{Recall} = \frac{TP}{TP+FN}$
 $\text{Precision} = \frac{TP}{TP+FP} \quad \text{FPR} = \frac{FP}{FP+TN}$
 $\text{FDR} = \frac{FP}{FP+TP}$
 $F_1 = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} = \frac{2TP}{2TP+FP+FN}$
ROC: FPR vs TPR=Recall
PR: Recall vs Precision